

AN ALTERNATIVE FOR THE IMPLEMENTATION OF KRUSKAL'S MINIMAL SPANNING TREE ALGORITHM

Jyrki KATAJAINEN and Olli NEVALAINEN

Department of Mathematical Sciences, University of Turku, Turku, Finland

Communicated by J. Nievergelt

Received August 1982

Revised December 1982

Abstract. An application of the bucket sort in Kruskal's minimal spanning tree algorithm is proposed. The modified algorithm is very fast if the edge costs are from a distribution which is close to uniform. This is due to the fact that the sorting phase then takes for an m edge graph an $O(m)$ average time. The $O(m \log m)$ worst case occurs when there is a strong peak in the distribution of the edge costs.

1. Introduction

Let $G = (N, E)$ be an undirected connected graph, N being the set of vertices and E the set of edges. Further let $E' \subseteq E$. Then the subgraph $T = (N, E')$ of G is a *spanning tree* of G if and only if T is a tree. Let us suppose that a cost or weight $c(e)$ is associated with each edge $e \in E$. The *minimal spanning tree* problem is then to find a spanning tree T , for which the sum of the costs

$$C(T) = \sum_{e \in T} c(e)$$

is minimal. In the following we denote $n = |N|$ (the number of vertices in the graph) and $m = |E|$ (the number of the edges).

The implementation of the well-known minimal spanning tree algorithm of Kruskal [7] is discussed in this paper. The impulse to this study arose from the work of Haymond, Jarvis and Shier [6] who gave an efficient implementation of Prim's algorithm. The so-called address calculation sort [9] is used in their algorithm to aid the selection of new candidate edges to be included in the spanning tree. The experimental tests are promising for the algorithm but it contains the limitation that the edge costs are (or must be mapped to) natural numbers from a limited range.

In the present paper we show how a distribution dependent sort method can be used also in Kruskal's algorithm. Instead of the address calculation sort we shall use the *hybrid sorting technique* discussed by Meijer and Akl [11]. The technique is more general than the former one and strong theoretical results on the average time complexity are available [1].

The overall plan of the paper is to give a description of the new algorithm in Section 2; to consider the time complexity of the average and worst cases in Section 3; and to give some experimental results in Section 4. Finally in Section 5 we shortly describe the application of hybrid sorting techniques in other graph problems.

2. A modified Kruskal's algorithm

Kruskal's minimal spanning tree algorithm uses the *greedy method* where the edges are considered in increasing order of the costs and included in the set T of the selected edges if the edges in T do not form a cycle also after the possible inclusion. The step of selecting an edge is repeated until $n - 1$ edges are included in T . Then T forms the wanted minimal spanning tree of G . The basic form of the above algorithm is the following [7]:

```

procedure Kruskal:
  {This program constructs the minimal spanning tree  $T$  for a connected
    $n$ -vertex graph  $G(N, E)$ }
  begin
     $T := \emptyset$ ;1
    while  $|T| < n - 1$  do begin
      Select an edge  $e'$  of lowest cost from  $E$ ;
      Delete  $e'$  from  $E$ ;
      if  $T \cup \{e'\}$  does not contain a cycle then  $T := T \cup \{e'\}$ ;
    end;
  end.
  
```

The choosing of an edge of lowest cost is usually accomplished by forming a min-heap of the edge costs. Thus the sorting of all edges can possibly be avoided. Determining whether the inclusion of a candidate edge would create a cycle in T can be seen as a UNION-FIND problem and it can be solved for example by using one of the possible tree structures as the data structure in the algorithm. One such is the algorithm which makes *Quick Merge* with *Weighting* rule in UNION and *Collapsing* rule in FIND (abbreviated hereafter QMWC) [7].²

¹ \emptyset denotes an empty data structure and also an empty set.

² The data structure used to represent a subset of vertices, forming a *connected component* of T , is a *rooted tree*, the *root* being a special vertex of the subset. Let $T(i)$ denote the tree that currently contains the vertex i . Let i and j be two vertices of N . The QMWC-algorithm performs UNION- and FIND-operations as follows:

FIND(i) (Determine the connected component containing the vertex i): Move along the father links to $root(T(i))$. After this apply the *collapsing rule*: if j is a node on the path from i to $root(T(i))$ then set $father(j) := root(T(i))$.

UNION(i, j) (Join the connected components containing the vertices i and j): Find $root(T(i))$ and $root(T(j))$. Then apply the *weighting rule*: if the number of nodes in $T(j)$ is less than the number in $T(i)$, make $root(T(i))$ the father of $root(T(j))$. Otherwise make $root(T(j))$ the father of $root(T(i))$.

Let us take another look at the search for the minimal edges. Denote by \min and \max the minimal and maximal edge costs and suppose that both are finite. We divide the range $[\min, \max]$ into b intervals of equal length and give the intervals the indexes $1, 2, \dots, b$. The edge e , the cost of which is $c(e)$, is associated with the bucket $j(e)$, where

$$j(e) = \left\lfloor \frac{c(e) - \min}{\max - \min} (b - 1) \right\rfloor + 1.$$

The edges belonging to the same interval j form a bucket denoted by $E(j)$. In the minimal spanning tree algorithm we first group the edges of E into buckets by the above formula. Thus instead of one large set E of the edges we now have a partitioning $E = E(1) \cup E(2) \cup \dots \cup E(b)$ of it. If the distribution of the edge costs is sufficiently close to uniform, the number of empty buckets is not large.

When choosing the minimal cost edges, the first non-empty bucket $E(j)$ among the buckets $E(1), E(2), \dots, E(b)$ is considered; a min-heap $H(j)$ of the edge costs is constructed for $E(j)$; minimal elements are removed from the heap and added to the spanning forest until the heap becomes empty or the minimal spanning tree T is ready. In the first case the next non-empty bucket is searched for and the same selection process is repeated.

The above gives a new form for Kruskal's algorithm:

procedure Our-Kruskal:

begin

 Determine the minimal and maximal costs of the edges in E ;

 Group the edges of E into buckets $E(1), E(2), \dots, E(b)$;

$T := \emptyset$;

 Put each vertex of N into a singular set;

$j := 0$; $H(j) := \emptyset$;

while $|T| < n - 1$ **do begin**

if $H(j) = \emptyset$ **then begin**

 Select the next j for which the bucket $E(j)$ is non-empty;

 Form a heap $H(j)$ for $E(j)$;

end;

 Select the minimal cost edge $e' = (u, v)$ from $H(j)$;

 Delete e' from $H(j)$;

if $\text{FIND}(u) \neq \text{FIND}(v)$ **then begin**

$T := T \cup \{e'\}$;

$\text{UNION}(\text{FIND}(u), \text{FIND}(v))$;

end;

end;

end.

A clearcut method to form the buckets is to link the elements, describing the edges, in the same bucket to form a linear list and to use an array of list heads

which point to the front of the lists. Note that this is not the only possible storage organization. One could for example count the edges in each bucket and then reorganize them according to the counts, cf. the method of Math Sort [5]. It is also possible to arrange the edges by making a chain of exchange operations, see [12].

There is more than one way of processing with the min-heaps. Firstly, when constructing a min-heap of a new bucket it is unnecessary to include into the bucket an edge which will cause a cycle in a connected component constructed so far. The appearance of a cycle can again be recognized by performing the FIND-operations for the two end points of the edge. In the following we suppose that the edges forming cycles are excluded at the moment of constructing the heaps.

Secondly, let us consider the operation of removing the smallest edge from the current heap. Here the edge appearing as the last element of the heap is moved to the place of the element just removed and then moved down to its right position in the heap. Kershenbaum and Van Slyke [8] considered a single heap-organization and proposed an improvement for this operation: when removing the last edge it is checked whether it would form a cycle in the current forest. In that case they delete the edge and consider the next last element. In our algorithm the use of buckets aims at the creation of several small heaps. Although the technique of Kershenbaum and Van Slyke is theoretically appealing our tests indicated that it increases the total overhead at least for uniformly distributed costs. Thus we have not implemented it in our algorithm. However, the algorithm is only ca. 4 percent faster than if the technique were in use. So, the technique of Kershenbaum and Van Slyke would be advantageous in the case where the distribution of the edge costs is very irregular.

3. Analysis

Let us first consider the *running time of min-heap operations* while constructing the heaps and removing the candidate edges. It may happen that almost all edges hit a single bucket and the last edge which will be added to T is the very largest one in E . Then there exists a very big heap from which all elements are to be removed. Thus in the *worst case* the time used for *selecting the edges* is $O(m \log m)$ which is the same as when sorting by heapsort. This is true also for the original Kruskal's algorithm.

The *expected running time* used for the heap operations depends on the distribution of the edge costs. It was shown in [11] that the bucket sorting technique works in $O(m)$ time for *uniformly distributed random numbers* and for a wide class of *smooth distributions*. (It is demanded that the distribution function $f(x)$ of the edge costs fulfils the conditions: $f(x) = 0$ for $x \notin [p, q]$ where q and p are fixed and $f(x)$ is finite. A still wider class of distributions is given in [3].) In the selection of the edges we essentially sort a subset of the edges by bucket sort and thus for these distributions the expected selection time is $O(m)$.

Let us next take a look at the *FIND*- and *UNION*-operations. We supposed that these are done by the QMWC-algorithm. It has been shown [13] that if $t(m', n)$ denotes the worst case time required to process an intermixed sequence of $m' \geq n$ FINDs and $n - 1$ UNIONs then

$$k_1 \cdot m' \cdot \alpha(m', n) \leq t(m', n) \leq k_2 \cdot m' \cdot \alpha(m', n)$$

for some positive constants k_1 and k_2 . Here $\alpha(m', n)$ is a very slowly increasing function for which $i \leq \alpha(m', n) \leq 3$ for all practical values of m' and n . Now we have $m' \leq 4m$ because the initialization of the heaps is preceded by at most $2m$ FINDs and possibly we have to remove all m edges from the heaps. Thus the *worst case* running time of the QMWC-algorithm in Kruskal's method is still almost linear on m .

The question on the *average running time* of the QMWC-algorithm in a minimal spanning tree algorithm seems to be difficult. The sequence of operations when introducing new edges to the spanning forest determines the growth of the subsets and if we want to determine the average running time we must fix the set of graphs for which we are solving the problem.

Yao analyzed in [16] several UNION-FIND algorithms in the case of the so-called *random spanning tree model*. He estimates the running time of a sequence of the equivalence operations " $i \equiv j$ ", which means an operation of the form

"if FIND(i) \neq FIND(j) then UNION(i, j)".

He defines the distribution of the input sequence by defining an ensemble Γ of instruction sequences and assuming that every sequence in Γ is equally likely to occur:

$\Gamma = \{(i_1 \equiv j_1, i_2 \equiv j_2, \dots, i_{n-1} \equiv j_{n-1}) \mid \text{the edges } (i_1, j_1), (i_2, j_2), \dots, (i_{n-1}, j_{n-1}) \text{ form a spanning tree on the vertices } \{1, 2, \dots, n\}\}.$

Then the QMW (the *Quick Merge with Weighting rule*) and thus also QMWC runs in an expected $O(n)$ time. But now if m , the number of edges in the graph, is $n - 1$ the spanning tree found in this way is also minimal. In the special case that additionally the edge costs are from a smooth distribution with short tails, the modified Kruskal's algorithm runs in an expected $O(m)$ time.³

³ Yao has also studied another model, called the *random graph model*, to make the equivalence operations. Here a sequence of distinct random edges are introduced into the graph consisting initially of n vertices and none edges. Knuth and Schönhage [10] have shown that for the QFW-algorithm (Quick Find with Weighting rule, where the sets are linearly linked lists) the average running time to do the UNION-operations until the graph is connected is $O(n)$. Thus the result can be interpreted as a spanning tree of a *complete graph* if the edges which form a cycle are rejected. On the other hand Yao has shown that the average running time of QFW and QMW is the same for this model when $m = n - 1$ but the case of a complete graph is not explicitly treated. Thus the average running time of the QMWC-algorithm remains undetermined for a general m .

A straightforward implementation of Kruskal's algorithm needs a *storage space* for the edges, costs and for the minimal spanning tree. For the QMWC-algorithm we need the size of each equivalence tree and the father links. The size can be stored as negative in the father link of the root. The total storage space is thus $2m + m + 2n + n = 3m + 3n$.

In our implementation of Kruskal's algorithm we have additionally to form the buckets. The running time depends on b , the number of buckets. The experiments (see Section 4) indicate that $b = m/10$ is a reasonable selection. Then we need $m/10$ list heads, an array of m links, and an auxiliary array of the length m for the heaps. The total storage space needed by the algorithm is thus about $5.1m + 3n$ storage locations.

4. Profiling the running time

To determine the actual running time of the new algorithm we used the model for generation of random graphs described by Haymond, Jarvis and Shier [6]. Here we first fix n and m and randomly connect the n vertices by $n - 1$ edges. This is done incrementally by considering a set $S(j) = \{1, 2, \dots, j\}$ of vertices and connecting the vertex $j + 1$ to a vertex k selected randomly from $S(j)$. Secondly, we add the remaining $m - (n - 1)$ edges to the connected graph with $n - 1$ edges randomly without repetitions. Finally, we associate costs drawn from a given distribution to the existing edges.

The two versions of Kruskal's algorithm were written in Pascal and all tests were performed in a DECSYSTEM-20 with a KL10 processor during a low period of the computer usage. The original Kruskal's algorithm was implemented in the form presented by Kershenbaum and Van Slyke [8].

As a first step, we wanted to test the effect of the parameter b , the number of buckets, on the running time of the modified Kruskal's algorithm. Table 1 shows for uniformly distributed edge costs the dependence of the observed average running time on the selection of b . It is observed that the running times do not change very

Table 1
The running times of our algorithm in milliseconds as a function of the number of buckets (b). The values are means of ten graphs ($n = 200$) with uniformly distributed edge costs

	<i>b</i>												
<i>m</i>	100	200	400	600	800	1000	1200	1600	2000	2400	4800	9600	14 400
1990	146	137	136	134	135	133	136	139					
9950		368	358	357	360	359	364	369	367	373	385	401	
17 910		627	576	563	563	566	569	571	585	587	595	616	627

radically when changing b . A selection $b = m/10$ seems to be reasonable and we use this in what follows. When sorting uniformly, normally or negative exponentially distributed random numbers the best selection of b was ca. 5 [12]. This result is in sound with the results of Table 1: in Kruskal's algorithm some of the edges are deleted already at the time when constructing a heap. Thus the heaps are in many cases smaller than the buckets from which they originate.

Table 2 shows the results of some test runs with uniformly distributed random costs. The values in Table 2 are means from ten repetitions. The results clearly indicate the power of bucket sort. In comparison to the straightforward version of Kruskal's algorithm the observed running time is reduced when using the new implementation and the reduction is the better the more edges we have in the graph. For the original Kruskal's algorithm the FIND-operations with QFW gave somewhat better running times than with QMWC. Note that the same observation was reported also by Haymond, Jarvis and Shier [6]. For our-Kruskal QFW and QMWC seem to give quite similar running times, see Table 2. (A statistical analysis of the differences was not performed.) The running times seem to be linear on m when our-Kruskal with QMWC is in use.

For the purpose of comparison we have also solved the same problems by the Whitney's publication of Prim's algorithm [15], see Table 2. This algorithm is written in FORTRAN and it should be noted that in the case of our compilers the FORTRAN versions of the same program tend to be somewhat faster than Pascal versions. Of the three spanning tree algorithms the Prim's algorithm was fastest for large m . Note that the two algorithms apply different organizations of the storage. In Kruskal's algorithm the graph is given in sequential list form whereas in Prim's algorithm the adjacency matrix is used for. The running time of Prim's algorithm is for a fixed n independent on m .

We made also some test runs with normally distributed edge costs. The graphs were exactly the same as before but costs were generated independently from a normal distribution with mean 0 and standard deviation 500. In the first set of test runs negative costs were allowed. Then the performance of our algorithm was almost as good as with uniformly distributed costs. In the second set of runs only positive costs were allowed i.e. we generated real numbers from the normal distribution mentioned above and take the absolute values of these numbers. Now the buckets with low indexes become much larger than the others. The performance of our algorithm was now somewhat weaker than in the uniform case, see Table 3.

Finally, we generated sparse connected graphs of grid type. The vertices of a graph are placed at random on the rectangular grid $\{(x, y) | x, y = 0, 1, 2, \dots, 999\}$. Edges are generated randomly as before but the cost of an edge is now the Euclidean distance between the end points. The results of test runs are shown in Table 4. The bucket version is also here much faster than a single heap version of Kruskal's algorithm.

300		299		4485	13 455				
	KRU(QFW)	137		719		1370			
	OUR(QFW)	109		315		557			
	OUR(QMWC)	84		238		499			
	WHI	464		464		462			
							465	461	461
350		349		6107	18 322				460
	KRU(QFW)	175		953		1766			
	OUR(QFW)	139		405					
	OUR(QMWC)	98		304					
	WHI								
400		399		7980					
	KRU(QFW)	197		1202					
	OUR(QFW)	155		510					
	OUR(QMWC)	112		383					
	WHI								
450		449		10 102					
	KRU(QFW)	250		1531					
	OUR(QFW)	200		679					
	OUR(QMWC)	126		467					
	WHI								
500		499		12 475					
	KRU(QFW)	283		1855					
	OUR(QFW)	225		791					
	OUR(QMWC)	141		582					
	WHI								

KRU = Kruskal's algorithm
 OUR = Kruskal's algorithm with bucket sort
 WHI = Whitney's publication [15] of Prim's algorithm
 QFW = Set operations made by Quick FIND with Weighting rule in the UNIONS
 QMWC = Set operations made by Quick Merge with Weighting rule in the UNIONS and Collapsing rule in the FINDs

Table 3

The observed average running time of Kruskal's algorithm with bucket sort in milliseconds for uniformly distributed and for normally distributed edge costs

	<i>m</i>					
	199	1990	5970	9950	13 930	17 910
<i>n</i> = 200						
Uniform OUR	56	131	234	339	449	588
Normal OUR	59	149	250	367	498	593

Table 4

The running times of Kruskal's algorithm (KRU) and Kruskal's algorithm with bucket sort (OUR) in milliseconds for sparse connected graphs ($m = 3n$) of grid type

	<i>n</i>									
	100	200	300	400	500	600	700	800	900	1000
KRU	80	201	348	540	724	976	1224	1536	1870	2219
OUR	50	97	147	194	253	309	363	411	465	525

5. Shortest route problems

The minimal spanning tree problem is not the only problem in graph theory gaining advantage of the bucket search methods. One such problem is that of determining the shortest routes from a given start point (point 1) to each other point of a directed network [2, 4]. One of the methods to solve this problem (Dijkstra's method) *reaches* new points thru the point which is on the minimal temporary distance from the point 1. This so-called *label setting method* sets the distance (label) of each point exactly once and needs some means to choose the point for which the path to point 1 is shortest. Also here a min-heap can be used for the selection. For a graph with edge costs attaining integers as their values Dial et al. [4] have applied the address calculation sort in the place of a heap, cf. [6]. The practical tests have revealed that the method is fast when the edge costs are from a reasonable range but it may waste space. For large random graphs with edge costs from a limited range the observed running times have been very promising but for graphs with a low number of edges the time is weaker than that of so-called *label correcting* techniques [4].

When the edge costs attain real numbers as their values Denardo and Fox [2] have used a bucked scheme resembling that of our paper. As noted above the node which is closest to the point 1 is selected for the point of the continuation in the label setting method. The labels can therefore be divided in a number of buckets

according to their values. This is much alike our approach but because the minimal and maximal values of the labels in the buckets (or in the heap) are gradually growing the buckets move 'to right'. This is achieved by making the bucket heads to form a circular list. This makes the management of the data structure more complicated than in the present paper where the top of the bucket area is static. On the other hand no cumbersome set operations are needed in the shortest route algorithm. Denardo and Fox have also generalized the same idea to a k -level bucket scheme for which they have given nice theoretical results of the worst case running time.

In practical tests [2] with a class of random grid networks with Euclidean distances the 1- and 2-level systems showed to be slower than a method based on label correcting. The opposite is true when the distances are taken from a truncated normal distribution. Because the selection of the nodes with the minimal distances greatly resembles the basic step in the bucket sort [12] one might expect that the 1-level organization is mainly of its simplicity the fastest in practice. Therefore it might be interesting to make a further study of the 1-level bucket system of Denardo and Fox and try to answer to questions like; what is the expected running time of the algorithm for a certain class of random networks; how to manage with a small amount of storage space; and how to update the buckets economically.

6. Concluding remarks

The use of bucket sort in Kruskal's minimal spanning tree algorithm was discussed. The motivation of the new method was to decrease the size of heaps by which the selection of the minimal cost edges is done. For graphs with uniformly distributed edge costs the new method is capable of reducing the running time of Kruskal's algorithm. On the other hand it is easy to construct a graph for which all except one of the edges fall into the same bucket and the method degenerates to the original one (with $O(m)$ extra work).

It was shown that in a very special case the expected running time is $O(m)$. The question is open for a general m . At least the experimental results give a reason to try to prove the linearity.

Weide [14] has proposed in a bucket sort use of an empirical cumulative distributive function to smoothen the distribution. We performed also some test runs with a bucket management algorithm of this kind but the observed running times suffered from the more complicated program code.

One can also apply the bucket sort technique in Prim's algorithm. The overall organization of the program is then the same as described by Haymond, Jarvis and Shier [6]. In the place of the synonym lists and the fixed mapping function we have a data structure similar to that in Section 2. Now the difficulty lies in the dynamic nature of the buckets. New edges may occur also in buckets which have already been bypassed. In addition because the number of edges in the bucket area

varies it is difficult to determine an appropriate value for the parameter b . We implemented also this variation and used in the tests the parameter value $b = n/2$. For very sparse graphs the bucket version of Prim's algorithm was the fastest of the algorithms studied in this paper whereas for dense graphs it gave very weak results.

References

- [1] S.G. Akl and H. Meijer, On the average-case complexity of 'bucketing' algorithms, *J. Algorithms* **3** (1982) 9–13.
- [2] E.V. Denardo and B.L. Fox, Shortest-route methods: 1, reaching, pruning and buckets, *Operations Res.* **27**(1) (1979) 161–186.
- [3] L. Devroye and T. Klinecsek, Average time behavior of distributive sorting algorithms, *Computing* **26**(1) (1981) 1–7.
- [4] R. Dial, F. Glover, D. Karney and D. Klingman, A computational analysis of alternative algorithms and labeling techniques for finding shortest path trees, *Networks* **9** (1979) 215–248.
- [5] W. Feurzeig, Algorithm 23: Math sort, Collected algorithms from ACM.
- [6] R.E. Haymond, J.P. Jarvis and D.R. Shier, Computational methods for minimum spanning tree problems, Technical Report # 354, Department of Mathematical Science, Clemson University (1980).
- [7] E. Horowitz and S. Sahni, *Fundamentals of Computer Algorithms* (Pitman, Potomac, MD, 1978).
- [8] A. Kershenbaum and R. Van Slyke, Computing minimum spanning trees efficiently, *Proc. 25th Annual Conference of the ACM* (1972) 518–527.
- [9] D.E. Knuth, *The Art of Computer Programming, Volume 1: Fundamental Algorithms* (Addison-Wesley, Reading, MA, 1973).
- [10] D.E. Knuth and A. Schönhage, The expected linearity of a simple equivalence algorithm, *Theoret. Comput. Sci.* **6** (1978) 281–315.
- [11] H. Meijer and S. G. Akl, The design and analysis of a new hybrid sorting algorithm, *Information Processing Lett.* **10**(4) (1980) 213–218.
- [12] O. Nevalainen and J. Ernvall, Implementation of a distributive sorting algorithm, Department of Mathematical Sciences, University of Turku, Finland (1981).
- [13] R. Tarjan, Efficiency of a good but not linear set union algorithm, *J. ACM* **22**(2) (1975) 215–225.
- [14] B.W. Weide, Statistical methods in algorithm design and analysis, Ph.D. Thesis, Carnegie-Mellon University (1978).
- [15] V.K.M. Whitney, Algorithm 422: Minimal spanning tree, *Comm. ACM* **15** (4) (1972) 273–274.
- [16] A.C. Yao, On the average behavior of set merging algorithms, *Proc. 8th Annual Symposium on Theory of Computing* (1976) 192–195.